



# X64 NASM Assembly Exercises: Algorithmic Programming

## Easy Exercises (1-2 hours each)

### String Operations

1. **strlen** - Implement `int strlen(const char* str)` that returns the length of a null-terminated C string.
2. **strcpy** - Implement `char* strcpy(char* dest, const char* src)` that copies a string from `src` to `dest` and returns `dest`.
3. **strchr** - Implement `char* strchr(const char* str, int c)` that finds the first occurrence of character `c` in string `str`.

### Math - Basic

4. **fibonacci** - Implement `long fib(int n)` that computes the  $n$ -th Fibonacci number ( $n \leq 40$ ).
5. **factorial** - Implement `long factorial(int n)` that computes  $n!$  for  $n \leq 20$ .
6. **gcd** - Implement `int gcd(int a, int b)` using Euclidean algorithm.
7. **power\_mod** - Implement `long power_mod(long base, long exp, long mod)` that computes  $(base^{exp}) \% mod$ .
8. **is\_prime** - Implement `int is_prime(long n)` that returns 1 if  $n$  is prime, 0 otherwise (trial division).

### Array Operations

9. **array\_sum** - Implement `long array_sum(const long* arr, int n)` that sums all elements in an array.
10. **array\_min\_max** - Implement `void array_min_max(const int* arr, int n, int* min, int* max)` that finds both min and max in one pass.
11. **array\_reverse** - Implement `void array_reverse(int* arr, int n)` that reverses an array in-place.
12. **memset** - Implement `void* memset(void* ptr, int value, size_t num)` that fills memory with a constant byte.

## Bit Operations

13. **popcount** - Implement `int popcount(unsigned long x)` that counts the number of set bits (without using POPCNT instruction).
14. **reverse\_bits** - Implement `unsigned int reverse_bits(unsigned int x)` that reverses the bit order in a 32-bit integer.
15. **is\_power\_of\_two** - Implement `int is_power_of_two(unsigned long x)` efficiently.

## CPU/System

16. **rdtsc\_read** - Implement `unsigned long read_tsc(void)` that reads the CPU timestamp counter using RDTSC.
17. **cpu\_brand** - Implement `void get_cpu_brand(char* buffer)` that retrieves CPU brand string using CPUID (buffer size  $\geq 48$  bytes).

## Computational Geometry - Basic

18. **point\_distance** - Implement `float point_distance(float x1, float y1, float x2, float y2)` computing Euclidean distance between two 2D points.
19. **point\_in\_rect** - Implement `int point_in_rect(float px, float py, float rx, float ry, float rw, float rh)` checking if point is inside rectangle.

## Hash/Checksum

20. **simple\_hash** - Implement `unsigned int simple_hash(const char* str)` using the djb2 hash algorithm.
21. **xor\_checksum** - Implement `unsigned char xor_checksum(const unsigned char* data, int len)` that XORs all bytes.

## Number Theory

22. **lcm** - Implement `long lcm(long a, long b)` that computes least common multiple.
23. **mod\_inverse** - Implement `long mod_inverse(long a, long m)` using extended Euclidean algorithm (assume  $\text{gcd}(a,m)=1$ ).

## Data Structures - Basic

24. **linear\_search** - Implement `int linear_search(const int* arr, int n, int target)` returning index or -1.
25. **binary\_search** - Implement `int binary_search(const int* sorted_arr, int n, int target)` on sorted array.

## Medium Exercises (3-5 hours each)

### String Operations - Optimized

1. **strcmp** - Implement `int strcmp(const char* s1, const char* s2)` with two versions: scalar and AVX2-optimized. Compare performance on strings of various lengths (10, 100, 1000, 10000 bytes).
2. **memcpy\_optimized** - Implement `void* memcpy(void* dest, const void* src, size_t n)` with AVX2 optimization for large blocks. Compare with scalar version.
3. **strstr** - Implement `char* strstr(const char* haystack, const char* needle)` using naive algorithm and compare with Boyer-Moore-Horspool variant.

### Math - Intermediate

4. **matrix\_multiply** - Implement `void matmul(const float* A, const float* B, float* C, int n)` for  $n \times n$  matrices using AVX/AVX2 instructions. Test with  $n = 8, 16, 32, 64$ .
5. **fibonacci\_matrix** - Implement `long fib_fast(int n)` using matrix exponentiation method for computing large Fibonacci numbers ( $n \leq 90$ ).
6. **sieve\_eratosthenes** - Implement `int sieve(int n, int* primes)` that finds all primes up to  $n$  using Sieve of Eratosthenes. Return count of primes.
7. **polynomial\_eval** - Implement `double poly_eval(const double* coeffs, int degree, double x)` using Horner's method with AVX2 for SIMD evaluation of 4 polynomials simultaneously.

### Computational Geometry

8. **rect\_intersection** - Implement `int rect_intersect(float x1, float y1, float w1, float h1, float x2, float y2, float w2, float h2, float* out_x, float* out_y, float* out_w, float* out_h)` that computes intersection of two axis-aligned rectangles. Return 1 if intersects, 0 otherwise.
9. **convex\_hull\_2d** - Implement `int convex_hull(const float* points_xy, int n, int* hull_indices)` using Graham scan algorithm for 2D points. Return number of hull vertices.

### Algorithms - Data Structures

10. **quicksort** - Implement `void quicksort(int* arr, int n)` using in-place partitioning.
11. **rolling\_hash** - Implement `unsigned long rolling_hash(const char* str, int window_size, unsigned long* hashes, int* count)` that computes all rolling hashes of given window size using Rabin-Karp polynomial rolling hash.

### Hash/Crypto

12. **crc32** - Implement `unsigned int crc32(const unsigned char* data, size_t len)` using table-lookup method (IEEE polynomial 0x04C11DB7).

## Graph Algorithms

13. **adjacency\_matrix\_transpose** - Implement `void transpose_graph(const int* adj, int* adj_t, int n)` that transposes an adjacency matrix representation of a directed graph.

## CPU Features

14. **cpuid\_features** - Implement `unsigned long get_cpu_features(void)` that detects and returns a bitmask of CPU features (SSE, AVX, AVX2, AVX512F, AES-NI, etc.) using CPUID.
15. **memcmp\_avx2** - Implement `int memcmp(const void* s1, const void* s2, size_t n)` with AVX2 optimization. Compare performance with scalar version on aligned/unaligned data.

## Hard Exercises (Multiple evenings)

### Advanced Math & SIMD

1. **arctan\_avx** - Implement `void arctan_avx(const float* input, float* output, int n)` computing arctangent using Taylor series or CORDIC algorithm with AVX/AVX512 instructions. Compare precision and performance with `atanf` from `libm` for  $n=1000000$  random inputs in  $[-1, 1]$ .
2. **matrix\_multiply\_avx512** - Implement `void matmul_avx512(const float* A, const float* B, float* C, int n)` using AVX512 instructions with loop tiling/blocking optimization. Compare with AVX2 version. Test with  $n = 128, 256, 512, 1024$ .
3. **fft\_complex** - Implement `void fft(float* real, float* imag, int n)` computing Fast Fourier Transform (Cooley-Tukey radix-2) for complex numbers.  $n$  must be power of 2. Test with  $n = 256, 1024$ .

### String/Encoding

4. **base64\_encode\_decode** - Implement both `int base64_encode(const unsigned char* input, int len, char* output)` and `int base64_decode(const char* input, unsigned char* output, int* out_len)`. Optimize encoding with SIMD. Test with random binary data and text files.
5. **regex\_simple** - Implement `int regex_match(const char* pattern, const char* text)` supporting only: literal characters, `.` (any char), `*` (zero or more of previous), `^` (start), `$` (end). Use backtracking algorithm.

### Computational Geometry - Advanced

6. **line\_segment\_intersection** - Implement `int segment_intersect(float x1, float y1, float x2, float y2, float x3, float y3, float x4, float y4, float* ix, float* iy)` that finds intersection point of two line segments using parametric equations. Return 1 if intersects, 0 if parallel/no intersection.
7. **point\_in\_polygon** - Implement `int point_in_polygon(float px, float py, const float* vertices_xy, int n)` using ray-casting algorithm to determine if point is inside polygon. Test with concave and convex polygons.

## Algorithms - Advanced

8. **floyd\_warshall** - Implement `void floyd_warshall(int* dist, int n)` computing all-pairs shortest paths. Input/output is  $n \times n$  adjacency matrix (INT\_MAX for no edge). Test with graphs of  $n = 50, 100, 200$  vertices.
9. **merge\_sort\_avx** - Implement `void mergesort(int* arr, int n)` with AVX2-optimized merge operation for sorting integers. Compare performance with scalar mergesort on arrays of size 10000, 100000, 1000000.

## Lock-Free Data Structures

10. **treiber\_stack** - Implement lock-free stack using Treiber's algorithm with `void push(struct node** head, long value)` and `int pop(struct node** head, long* value)` using CMPXCHG (compare-and-swap). Test with single-threaded correctness (push/pop 10000 elements). Note: focus on correct assembly implementation of CAS loop, actual concurrency testing optional.

## Testing Notes

### Test Data Sources:

- **String operations:** Generate random ASCII strings, use Lorem Ipsum text, Project Gutenberg texts
- **Math:** Use known sequences (OEIS), WolframAlpha for verification
- **Matrices:** Generate random matrices, use identity/zero matrices for edge cases
- **Geometry:** Generate random points/rectangles, use known geometric configurations
- **Graphs:** Use adjacency matrices from standard graph datasets (SNAP, NetworkX generated graphs)
- **Numerical:** Compare against `libm` functions, use known mathematical constants

### Performance Testing:

- Use RDTSC for cycle-accurate timing
- Repeat tests 100-1000 times and take median
- Test on cold and warm cache
- Use `perf stat` for hardware counter analysis

### AVX512 Note:

Your AMD CPU might not have AVX512. Check with:

```
grep avx512 /proc/cpuinfo
```

If not available, substitute AVX512 exercises with AVX2 equivalents or advanced AVX2 optimizations (e.g., better cache utilization, loop unrolling).

**Calling Convention:**

All functions follow System V AMD64 ABI (Linux x64):

- Integer/pointer args: RDI, RSI, RDX, RCX, R8, R9
- Float args: XMM0-XMM7
- Return: RAX (integer), XMM0 (float)
- Caller-saved: RAX, RCX, RDX, RSI, RDI, R8-R11, XMM0-XMM15
- Callee-saved: RBX, RBP, R12-R15

Good luck with your assembly programming practice!